

solve this problem [5/5]

Andrea Giugliano

January 14, 2018

Contents

1	the exercise	2
2	DONE the clojure setup	4
3	DONE the db setup	4
4	DONE the xml setup	5
4.1	DONE let's run the stackoverflow example as it is	6
4.2	DONE now let's play a bit to understand the basics	7
4.3	DONE let's understand this better	7
4.4	DONE let's use data.xml to have a lazy reading of xml	7
5	DONE test environment [4/4]	8
5.1	DONE create a test db	8
5.2	DONE try postgres-async	10
5.3	DONE create a test xml	11
5.3.1	DONE now we can test :)	11
6	DONE finally some code (by using Test Driven Development)	12
	[9/9]	
6.1	DONE check in the given sql file how they add records	12
6.2	DONE feature 0: create a clj representation of person record from xml	12
6.3	DONE feature 1: create a person record from xml	14
6.4	DONE feature 2: do not create a person record if exists in the db	17
6.5	DONE feature 3: update a person number if it exists in the db	22

6.6	DONE feature 4: update all the people in the xml returning overall stats	22
6.7	DONE refactor code further ("Write clean code")	27
6.8	DONE "minimize the overall run time of the merge process" .	29
6.9	DONE check performance (" Reason about performance and memory usage")	30

- CLOSING NOTE *[2017-01-22 Sun 20:35]*

See: Re: Application to Clojure Developer role

1 the exercise

Get the file:

```
cd /tmp
wget http://exchange.somecoolcompany.com.s3-website.eu-central-1.amazonaws.com/somecoolcompany-backend-test.tar
```

Untar the main file:

```
tar xvf somecoolcompany-backend-test.tar
```

```
# Somecoolcompany backend developer test
```

```
## The task
```

There is a PostgreSQL table of persons (person), uniquely identified by their first name (fname), last name (lname) and date of birth (dob). Every person has a telephone number (phone).

This table needs to be updated from an XML file containing elements of the form

```
'''
<member>
  <first-name>JOHN</first-name>
  <last-name>DOE</last-name>
  <date-of-birth>2002-02-01</date-of-birth>
  <phone>9548938821</phone>
</member>
'''
```

If the phone number is already correct, nothing should be changed in the database. If a person record does not exist, it needs to be created.

The person database table contains 10 million rows.

The update file contains 1.5 million entries.

Objective

- Write clean code that performs the operation correctly
- Provide basic loading statistics at the end of the operation
- Use proper mechanisms to process the input file
- Find ways to minimize the overall run time of the merge process
- Reason about performance and memory usage

The number of records in the sample database and the input file are meant to reflect the number of records in a production system. A production system would have more individual fields per person, consider that when choosing an implementation strategy.

Files

The file `person.sql.gz` contains a database dump of the person table which can be imported into PostgreSQL.

The file `update-file.xml.gz` contains the XML input file to be merged into the database.

Here Re: Application to Clojure Developer role Volker says that I can use open source libraries, and no deadline. Also about this:

The statistics should be useful to the user of the software and relate to the input business objects, not the database system.

I think they want to know just how many attributes were updated, how many were present already, and how many were created.

2 DONE the clojure setup

- CLOSING NOTE *[2017-01-18 Wed 17:08]*
- Note taken on *[2017-01-18 Wed 09:57]*
<https://clojuredocs.org/clojure.xml/parse>
- Note taken on *[2017-01-18 Wed 09:56]*
clj+postgres: <https://github.com/alaisi/postgres.async>

We create a new clojure project for the somecoolcompany-exercise:

```
lein new trysomecoolcompany-exercise
```

Then we modify the project metadata to include the somecoolcompany-exercise library:

```
(defproject trysomecoolcompany-exercise "0.1.0-SNAPSHOT"
  :description "somecoolcompany-exercise"
  :url "http://example.com/FIXME"
  :license {:name "Eclipse Public License"
            :url "http://www.eclipse.org/legal/epl-v10.html"}
  :dependencies [[org.clojure/clojure "1.8.0"]
                 [org.clojure/data.zip "0.1.2"]
                 [org.clojure/data.xml "0.2.0-alpha2"]
                 [alaisi/postgres.async "0.8.0"]]
  :main trysomecoolcompany-exercise.core)
```

Then we install the dependencies:

```
lein deps
```

3 DONE the db setup

<http://stackoverflow.com/questions/6842393/import-sql-dump-into-postgresql-database>

<https://www.postgresql.org/docs/current/static/role-attributes.html>

Now let's work with the database.

Unzip sql file:

```
gunzip -q person.sql.gz
```

I need to install and initialize postgres:

```
nix -iA nixos.postgresql
```

Then I have to initialize it, so let's create a directory for this:

```
mkdir db
initdb -D db
pg_ctl -D db -l logfile start
# we need to create the user hans
echo "CREATE ROLE hans LOGIN;"> patch.sql
# we need to alter the table to have large text fields
psql postgres < patch.sql
```

finally we import the provided sql dump after patching the NULL string with the dummy value 5500-12-12 (there are few NULL strings in the data values and they would make the primary key constraint invalid):

```
sed -i "s/\\N/5500-12-12/g" /tmp/somecoolcompany-backend-test/person.sql
psql postgres < /tmp/somecoolcompany-backend-test/person.sql
```

finally we add a primary key, or better unique indexes to avoid the NULL problem:

```
echo "ALTER TABLE person ADD PRIMARY KEY (fname, lname, dob);" > patch2.sql
#echo "CREATE UNIQUE INDEX p_col_uni_idx ON person (fname, lname, dob) WHERE fname IS NOT NULL;"
#echo "CREATE UNIQUE INDEX p_col_a_uni_idx ON person (fname, lname) WHERE dob IS NULL;"
#echo "CREATE UNIQUE INDEX p_col_b_uni_idx ON person (fname, dob) WHERE lname IS NULL;"
psql postgres < patch2.sql
```

4 DONE the xml setup

- CLOSING NOTE *[2017-01-18 Wed 18:18]*
I moved to data.xml to have lazy reading of the big xml file
- CLOSING NOTE *[2017-01-18 Wed 11:42]*

Unzip xml file:

```
gunzip -q update-file.xml.gz
```

<http://stackoverflow.com/questions/11537923/searching-xml-in-clojure>
there is also an xpath library, but maybe I do not need to query the file so much: <https://github.com/kyleburton/clj-xpath/>

Let's test a bit this zip library.

4.1 DONE let's run the stackoverflow example as it is

- CLOSING NOTE *[2017-01-18 Wed 11:18]*

<http://stackoverflow.com/questions/11537923/searching-xml-in-clojure>

The xml file:

```
<data>
  <products>
    <product>
      <section>Red Section</section>
      <images>
        <image>img.jpg</image>
        <image>img2.jpg</image>
      </images>
    </product>
    <product>
      <section>Blue Section</section>
      <images>
        <image>img.jpg</image>
        <image>img3.jpg</image>
      </images>
    </product>
    <product>
      <section>Green Section</section>
      <images>
        <image>img.jpg</image>
        <image>img2.jpg</image>
      </images>
    </product>
  </products>
</data>
```

```
(ns core
  (:use clojure.data.zip.xml)
  (:require [clojure.zip :as zip]
             [clojure.xml :as xml]))
```

```
(def data (zip/xml-zip (xml/parse "/tmp/test-xml.xml")))
(def products (xml-> data :products :product))
```

```
(for [product products :let [image (xml-> product :images :image)]
      :when (some (text= "img2.jpg") image)]
  {:section (xml1-> product :section text)
   :images (map text image)})
```

4.2 DONE now let's play a bit to understand the basics

- CLOSING NOTE *[2017-01-18 Wed 11:39]*
cool!

```
(ns core
  (:use clojure.data.zip.xml)
  (:require [clojure.zip :as zip]
             [clojure.xml :as xml]))
```

```
(def data (zip/xml-zip (xml/parse "/tmp/test-xml.xml")))
(def products (xml-> data :products :product))
```

```
(for [product products :let [image (xml-> product :images :image)]
      ; let's filter the Blue Section only
      :when (some (text= "Blue Section") (xml-> product :section))] ;
  {:section (xml1-> product :section text)
   :images (map text image)})
```

4.3 DONE let's understand this better

- CLOSING NOTE *[2017-01-18 Wed 11:42]*

<https://ravi.pckl.me/short/functional-xml-editing-using-zippers-in-clojure/>

So, now I recollect: we use the zipper data structure that was firstly published in a Haskell paper. Basically is an iterator that maintain its context so can iterate in whatever dimension.

The `xml->` and `xml1->` differ in the number of results returned.

4.4 DONE let's use data.xml to have a lazy reading of xml

- CLOSING NOTE *[2017-01-18 Wed 18:18]*

since we have a huge xml, we cannot have it all as a string.

clojure.data.xml to the rescue: it imports an element at a time from a reader.

However I should still be able to use the zipper to query it:

```
(ns core
  (:use clojure.data.zip.xml)
  (:require [clojure.zip :as zip]
            [clojure.xml :as xml]
            [clojure.data.xml :as data.xml]))

(def data (zip/xml-zip (data.xml/parse (clojure.java.io/input-stream "/tmp/test-xml.xml"))))

(def products (xml-> data :products :product))

(for [product products :let [image (xml-> product :images :image)]
      :when (some (text= "img2.jpg") image)]
  {:section (xml1-> product :section text)
   :images (map text image)})

[org.clojure/data.xml "0.0.8"]
```

5 DONE test environment [4/4]

- CLOSING NOTE *[2017-01-18 Wed 17:00]*

5.1 DONE create a test db

- CLOSING NOTE *[2017-01-18 Wed 16:12]*

Okay, I want to create a new table called `test_person` with some data in it:

```
--
-- PostgreSQL database dump
--

SET statement_timeout = 0;
SET lock_timeout = 0;
SET client_encoding = 'UTF8';
SET standard_conforming_strings = on;
SET check_function_bodies = false;
```



```

SET client_min_messages = warning;

--
-- Name: plpgsql; Type: EXTENSION; Schema: -; Owner:
--

CREATE EXTENSION IF NOT EXISTS plpgsql WITH SCHEMA pg_catalog;

--
-- Name: EXTENSION plpgsql; Type: COMMENT; Schema: -; Owner:
--

COMMENT ON EXTENSION plpgsql IS 'PL/pgSQL procedural language';

SET search_path = public, pg_catalog;

SET default_tablespace = '';

SET default_with_oids = false;

--
-- Name: person; Type: TABLE; Schema: public; Owner: hans; Tablespace:
--

CREATE TABLE test_person (
    fname character varying,
    lname character varying,
    dob date,
    phone character(10),
    PRIMARY KEY(fname, lname, dob) -- we need a primary key on the fields
);

ALTER TABLE person OWNER TO hans;

--
-- Data for Name: test_person; Type: TABLE DATA; Schema: public; Owner: hans
--

```

```
COPY test_person (fname, lname, dob, phone) FROM stdin;
JIARA HERTZEL 1935-06-05 5859012134
RONJARVIOU COMELLO 1932-09-27 7702713416
\.
```

```
--
-- Name: public; Type: ACL; Schema: -; Owner: hans
--
```

```
REVOKE ALL ON SCHEMA public FROM PUBLIC;
REVOKE ALL ON SCHEMA public FROM hans;
GRANT ALL ON SCHEMA public TO hans;
GRANT ALL ON SCHEMA public TO PUBLIC;
```

```
--
-- PostgreSQL database dump complete
--
```

and we import it in the database:

```
psql postgres < /tmp/trysomecoolcompany-exercise/resources/test_person.sql
```

5.2 DONE try postgres-async

- CLOSING NOTE *[2017-01-18 Wed 16:37]*

```
(require '[clojure.core.async :as async :refer [<! >! <!! timeout chan alt! go]])
(require '[postgres.async :refer :all])
```

```
(def db (open-db {:hostname "localhost"
                  :port 5432 ; default
                  :database "postgres"
                  :username "andrea"
                  :password ""
                  :pool-size 25})) ; default
```

```
;; let's get attributes of the test table
(<!! (query! db ["select * from test_person"])))
```

```
;; let's insert a test record
```

```

(<!! (insert! db {:table "test_person"} {:fname "Fiona", :lname "Lullaby", :dob "2010-02-01"}))

;; let's get the test row
(<!! (query! db ["select * from test_person where fname='Fiona'"])))

;; finally let's remove the test record
(<!! (execute! db ["delete from test_person where fname='Fiona' and lname='Lullaby' and dob='2010-02-01'"])))

;; now we cannot find the test row anymore
(<!! (query! db ["select * from test_person where fname='Fiona'"])))

(close-db! db)

```

5.3 DONE create a test xml

- CLOSING NOTE *[2017-01-18 Wed 09:55]*

```

<members>
  <member>
    <firstname>JOHN</firstname>
    <lastname>DOE</lastname>
    <date-of-birth>2002-02-01</date-of-birth>
    <phone>9548938821</phone>
  </member>
  <member>
    <firstname>Fiona</firstname>
    <lastname>Lullaby</lastname>
    <date-of-birth>2010-02-01</date-of-birth>
    <phone>9548938822</phone>
  </member>
  <member>
    <firstname>XX</firstname>
    <lastname>YY</lastname>
    <date-of-birth>2000-02-01</date-of-birth>
    <phone>9548138821</phone>
  </member>
</members>

```

5.3.1 DONE now we can test :)

- CLOSING NOTE *[2017-01-18 Wed 17:00]*

- Note taken on *[2017-01-18 Wed 11:37]*
to import the xml resource we can use

```
(ns core
  (:use clojure.data.zip.xml)
  (:require [clojure.zip :as zip]
             [clojure.xml :as xml]))
```

```
(def test-xml
  (-> "test.xml"
      clojure.java.io/resource
      clojure.java.io/input-stream
      xml/parse
      zip/xml-zip))
```

```
;; let's get the name for fun
(xml-> test-xml :members :member :first-name text)
```

```
nil#'core/test-xml("JOHN")
```

6 DONE finally some code (by using Test Driven Development) [9/9]

- CLOSING NOTE *[2017-01-22 Sun 20:27]*

6.1 DONE check in the given sql file how they add records

- CLOSING NOTE *[2017-01-18 Wed 17:04]*
no, they do not set any primary key on the table person
- Note taken on *[2017-01-18 Wed 11:11]*
fname,sname,date is the key tuple: they should add this constraint somehow

6.2 DONE feature 0: create a clj representation of person record from xml

- Note taken on *[2017-01-18 Wed 11:51]*
test that the representation can be converted back to the same xml

- keep in mind that the DB could own more fields

A production system would have more individual fields per person, consider that when choosing an implementation strategy.

First tests

```
(ns trysomecoolcompany-exercise.core-test
  (:require [clojure.test :refer :all]
             [trysomecoolcompany-exercise.core :refer :all]))

(def url
  (-> "test.xml"
      clojure.java.io/resource
      str))

(deftest xml->clj-vector-maps-test
  (testing "xml->clj should produce a list of maps with keys :fname :lname :dob :phone"
    (let [expected (list {:fname "JOHN"
                          :lname "DOE"
                          :dob "2002-02-01"
                          :phone "9548938821"
                          }
                          {:fname "Fiona"
                          :lname "Lullaby"
                          :dob "2010-02-01"
                          :phone "9548938822"
                          }
                          {:fname "XX"
                          :lname "YY"
                          :dob "2000-02-01"
                          :phone "9548138821"
                          })]
      (is (= (xml->clj url) expected))))))

(deftest xml->clj-keys-test
  (testing "xml->clj->db should contain only :fname :lname :dob :phone"
    (is (= (set (mapcat keys (xml->clj url))) #{:fname :lname :dob :phone}))))
```

Now code to pass them:

```

(ns trysomecoolcompany-exercise.core
  (:require [clojure.zip :as zip]
             [clojure.data.zip.xml :refer :all]
             [clojure.data.xml :as data.xml]))

(defn get-zipper
  "creates a zipper from an xml url"
  [url]
  (-> url
      clojure.java.io/input-stream
      data.xml/parse
      zip/xml-zip))

(defn xml->clj
  "transforms zipper nodes in maps"
  [url]
  (let [zxml (get-zipper url)
        members (xml-> zxml :members :member)]
    (for [member members :let [fname (xml1-> member :first-name text)
                                lname (xml1-> member :last-name text)
                                dob (xml1-> member :date-of-birth text)
                                phone (xml1-> member :phone text)]]
      ]
      {:fname fname :lname lname :dob dob :phone phone})))

```

6.3 DONE feature 1: create a person record from xml

- CLOSING NOTE *[2017-01-19 Thu 00:42]*

Let's add the database now:

```

(ns trysomecoolcompany-exercise.core-test
  (:require [clojure.test :refer :all]
             [trysomecoolcompany-exercise.core :refer :all]))

(def db-conf {:host "localhost"
              :port 5432
              :dbname "postgres"
              :user "andrea"
              :password ""
              :table "test_person"})

```

```

(def url
  (-> "test.xml"
      clojure.java.io/resource
      str))

(deftest xml->clj-vector-maps-test
  (testing "xml->clj should produce a list of maps with keys :fname :lname :dob :phone"
    (let [expected (list {:fname "JOHN"
                          :lname "DOE"
                          :dob "2002-02-01"
                          :phone "9548938821"
                          }
                          {:fname "Fiona"
                          :lname "Lullaby"
                          :dob "2010-02-01"
                          :phone "9548938822"
                          }
                          {:fname "XX"
                          :lname "YY"
                          :dob "2000-02-01"
                          :phone "9548138821"
                          })]
      (is (= (xml->clj url) expected))))))

(deftest xml->clj-keys-test
  (testing "xml->clj->db should contain only :fname :lname :dob :phone"
    (is (= (set (mapcat keys (xml->clj url))) #{:fname :lname :dob :phone}))))

(deftest clj-add-person-db-test
  (testing "xml-add-person-db should add a new person in the database from a person map"
    (let [db (connect-db db-conf)
          t (db-conf :table)
          person {:fname "XX"
                  :lname "YY"
                  :dob "2000-02-01"
                  :phone "9548138821"
                  }
          expected {:updated 1 :rows []}]
      (is (= (clj-add-person-db db t person) expected)))

```

```
(clj-del-person-db db t person)))) ;cleaning db after usage
```

Now code to pass them:

```
(ns trysomecoolcompany-exercise.core
  (:require [clojure.zip :as zip]
             [clojure.data.zip.xml :refer :all]
             [clojure.data.xml :as data.xml]
             [clojure.core.async :as async :refer [<!!]]
             [postgres.async :refer :all]))

(defn get-zipper
  "creates a zipper from an xml url"
  [url]
  (-> url
      clojure.java.io/input-stream
      data.xml/parse
      zip/xml-zip))

(defn xml->clj
  "transforms zipper nodes in maps"
  [url]
  (let [zxml (get-zipper url)
        members (xml-> zxml :members :member)]
    (for [member members :let [fname (xml1-> member :first-name text)
                                lname (xml1-> member :last-name text)
                                dob (xml1-> member :date-of-birth text)
                                phone (xml1-> member :phone text)]
          ]
      {:fname fname :lname lname :dob dob :phone phone})))

(defn connect-db
  "creates a db connection"
  [dbc]
  (open-db {:hostname (dbc :host)
            :port (dbc :port)
            :database (dbc :dbname)
            :username (dbc :user)
            :password (dbc :password)}
```



```

:pool-size 25}))

(defn close-db
  "closes a db connection"
  [db]
  (close-db! db))

(defn clj-add-person-db
  "add or update person to database's table"
  [db t p]
  (<!! (insert! db {:table t} p)))

(defn clj-del-person-db
  "delete person from database's table"
  [db t p]
  (<!! (execute! db [(str "delete from " t " where fname='"(p :fname)'" and lname='"(p

```

6.4 DONE feature 2: do not create a person record if exists in the db

- CLOSING NOTE *[2017-01-19 Thu 12:34]*
 woah, this was a little of postgres study: I like upsert (on conflict do set...)!

here the best choice is to add a primary key on name lastname and dob. However, we probably want to use an upsert, because using two postgres commands allows race conditions: <http://stackoverflow.com/questions/1109061/insert-on-duplicate-update-in-postgresql/8702291#8702291>

Let's add a test case for upsert:

```

(ns tryosomecoolcompany-exercise.core-test
  (:require [clojure.test :refer :all]
             [tryosomecoolcompany-exercise.core :refer :all]))

(def db-conf {:host "localhost"
              :port 5432
              :dbname "postgres"
              :user "andrea"
              :password ""
              :table "test_person"})

```

```

(def url
  (-> "test.xml"
      clojure.java.io/resource
      str))

(deftest xml->clj-vector-maps-test
  (testing "xml->clj should produce a list of maps with keys :fname :lname :dob :phone"
    (let [expected (list {:fname "JOHN"
                          :lname "DOE"
                          :dob "2002-02-01"
                          :phone "9548938821"
                          }
                          {:fname "Fiona"
                          :lname "Lullaby"
                          :dob "2010-02-01"
                          :phone "9548938822"
                          }
                          {:fname "XX"
                          :lname "YY"
                          :dob "2000-02-01"
                          :phone "9548138821"
                          })]
      (is (= (xml->clj url) expected))))))

(deftest xml->clj-keys-test
  (testing "xml->clj->db should contain only :fname :lname :dob :phone"
    (is (= (set (mapcat keys (xml->clj url))) #{:fname :lname :dob :phone}))))

(deftest clj-add-person-db-test
  (testing "clj-add-person-db should add a new person in the database from a person map"
    (let [db (connect-db db-conf)
          t (db-conf :table)
          person {:fname "XX"
                  :lname "YY"
                  :dob "2000-02-01"
                  :phone "9548138821"
                  }
          expected {:updated 1 :ignored 0}]
      (clj-del-person-db db t person) ; cleaning db before usage
      (is (= (clj-add-person-db db t person) expected)))

```

```

        (clj-del-person-db db t person)))) ;cleaning db after usage

(deftest clj-add-same-person-twice-db-test
  (testing "clj-add-people-db should add a new person in the database from a person map"
    (let [db (connect-db db-conf)
          t (db-conf :table)
          person {:fname "X"
                  :lname "Y"
                  :dob "1999-11-01"
                  :phone "9128138821"
                  }
          expected {:updated 1 :ignored 1}]
      (clj-del-person-db db t person) ; cleaning db before usage
      (is (= (clj-add-people-db db t person person) expected))
      (clj-del-person-db db t person)))) ;cleaning db after usage

(deftest clj-add-two-people-db-test
  (testing "clj-add-people-db should add two new person in the database from a person map"
    (let [db (connect-db db-conf)
          t (db-conf :table)
          p1 {:fname "ZX"
              :lname "Y"
              :dob "1899-11-01"
              :phone "9128138821"
              }
          p2 {:fname "FX"
              :lname "Y"
              :dob "1129-11-01"
              :phone "9128138821"
              }
          expected {:updated 2 :ignored 0}]
      (clj-del-person-db db t p1) ; cleaning db before usage
      (clj-del-person-db db t p2)
      (is (= (clj-add-people-db db t p1 p2) expected))
      (clj-del-person-db db t p1) ; cleaning db after usage
      (clj-del-person-db db t p2))))

```

Now code to pass the test:

```
(ns trysomecoolcompany-exercise.core
```

```

(:require [clojure.zip :as zip]
          [clojure.data.zip.xml :refer :all]
          [clojure.data.xml :as data.xml]
          [clojure.core.async :as async :refer [<!!]]
          [postgres.async :refer :all]))

(defn get-zipper
  "creates a zipper from an xml url"
  [url]
  (-> url
      clojure.java.io/input-stream
      data.xml/parse
      zip/xml-zip))

(defn xml->clj
  "transforms zipper nodes in maps"
  [url]
  (let [zxml (get-zipper url)
        members (xml-> zxml :members :member)]
    (for [member members :let [fname (xml1-> member :first-name text)
                                lname (xml1-> member :last-name text)
                                dob (xml1-> member :date-of-birth text)
                                phone (xml1-> member :phone text)]]
      ]
      {:fname fname :lname lname :dob dob :phone phone})))

(defn connect-db
  "creates a db connection"
  [dbc]
  (open-db {:hostname (dbc :host)
            :port (dbc :port)
            :database (dbc :dbname)
            :username (dbc :user)
            :password (dbc :password)
            :pool-size 25}))

(defn close-db
  "closes a db connection"
  [db]
  (close-db! db))

```

```

(defn clj-add-person-db
  "add or update person to database's table"
  [db t p]
  ; I needed to add a primary key over fname lname and dob to make
  ; this work
  (let [ks (keys p)
        kks (drop-last ks)
        pk (name (last ks))
        kcols (clojure.string/join " , " (map name kks))
        cols (clojure.string/join " , " (map name ks))
        vs (vals p)
        vals (str "'" (clojure.string/join "' , '" vs) "'") ;FIXME adding single quote
        pv (str "'" (last vs) "'")
        query (str
                "insert into " t
                " (" cols ") "
                "values (" vals ") "
                "on conflict (" kcols ") "
                "do update set " pk " = " pv "where "pv" <> " t "." pk ";"
                )
        (if (= {:updated 0 :rows []} (<!! (execute! db [query])))
            {:updated 1 :ignored 0}
            {:updated 0 :ignored 1})
        (reduce f {:updated 0 :ignored 0} rs)))

(defn clj-add-people-db
  "add or update person to database's table"
  [db t & ps]
  (let [rs (map #(clj-add-person-db db t %) ps)
        f (fn [acc r] {:updated (+ (r :updated) (acc :updated))
                        :ignored (+ (r :ignored) (acc :ignored))})]
    (reduce f {:updated 0 :ignored 0} rs)))

(defn clj-del-person-db
  "delete person from database's table"
  [db t p]
  (<!! (execute! db [(str "delete from " t " where fname='"(p :fname)'" and lname='"(p

```

6.5 DONE feature 3: update a person number if it exists in the db

- CLOSING NOTE *[2017-01-19 Thu 00:43]*
done during feature 2

6.6 DONE feature 4: update all the people in the xml returning overall stats

- CLOSING NOTE *[2017-01-19 Thu 14:14]*

Let's add a test case for a function that transfers xml people to db people:

```
(ns trysomecoolcompany-exercise.core-test
  (:require [clojure.test :refer :all]
             [trysomecoolcompany-exercise.core :refer :all]))

(def db-conf {:host "localhost"
              :port 5432
              :dbname "postgres"
              :user "andrea" ;; FIXME change with a postgres user
              :password ""})

(def db (connect-db db-conf))

(def t "test_person")

(def url
  (-> "test.xml"
      clojure.java.io/resource
      str))

(def p123
  [{:fname "JOHN"
    :lname "DOE"
    :dob "2002-02-01"
    :phone "9548938821"
    }
   {:fname "Fiona"
    :lname "Lullaby"
    :dob "2010-02-01"
    :phone "9548938822"
   }])
```

```

    }
    {:fname "XX"
     :lname "YY"
     :dob "2000-02-01"
     :phone "9548138821"
    })

(def p {:fname "XX"
       :lname "YY"
       :dob "2001-02-01"
       :phone "9548138821"
      })

(def person {:fname "X"
            :lname "Y"
            :dob "1919-11-01"
            :phone "9128138821"
           })

(def p1 {:fname "ZX"
        :lname "Y"
        :dob "1899-11-01"
        :phone "9128138821"
       })

(def p2 {:fname "FX"
        :lname "Y"
        :dob "1129-11-01"
        :phone "9128138821"
       })

(defn cool-down []
  (doall (map #(clj-del-person-db db t %) (conj p123 person p p1 p2)))
  (Thread/sleep 100)
  (close-db db))

; to clean database
(use-fixtures :once
  (fn [tests]
    (tests)

```

```

(cool-down)))

(deftest xml->clj-vector-maps-test
  (testing "xml->clj should produce a list of maps with keys :fname :lname :dob :phone"
    (let [expected p123]
      (is (= (into [] (xml->clj url)) expected)))))

(deftest xml->clj-keys-test
  (testing "xml->clj->db should contain only :fname :lname :dob :phone"
    (is (= (set (mapcat keys (xml->clj url))) #{:fname :lname :dob :phone}))))

(deftest clj-add-person-db-test
  (testing "clj-add-person-db should add a new person in the database from a person map"
    (let [expected {:updated 1 :ignored 0}]
      (is (= (clj-add-person-db db t p) expected)))))

(deftest clj-add-same-person-twice-db-test
  (testing "clj-add-people-db should add a new person in the database from two same people"
    (let [expected {:updated 1 :ignored 1}]
      (is (= (clj-add-people-db db t [person person]) expected)))))

(deftest clj-add-two-people-db-test
  (testing "clj-add-people-db should add two new people in the database from two people"
    (let [expected {:updated 2 :ignored 0}]
      (is (= (clj-add-people-db db t [p1 p2]) expected)))))

(deftest xml-db-transfer-people-test
  (testing "xml-db-transfer-people should insert not existing people in the db"
    (let [expected {:updated 3 :ignored 0}]
      (is (= (xml-db-transfer-people db t url) expected)))))

```

Now code to pass the test:

```

(ns trysofcoolcompany-exercise.core
  (:require [clojure.zip :as zip]
             [clojure.data.zip.xml :refer :all]
             [clojure.data.xml :refer [parse]]
             [clojure.data.xml :as data.xml]
             [clojure.core.async :as async :refer [<!!]]
             [postgres.async :refer :all])

```



```

[clojure.string :refer [join]])

(defn xml->clj
  "transforms zipper nodes in maps"
  [url]
  (let [get-members (fn [x] (->> x
                                clojure.java.io/input-stream
                                parse
                                :content
                                (filter #(= :member (:tag %)))))
        get-value (fn [key] (xml1-> (zip/xml-zip member) key text))]
    (for [member (get-members url)] :let [fname (get-value :firstname)
                                           lname (get-value :lastname)
                                           dob   (get-value :date-of-birth)
                                           phone (get-value :phone)]
      ]
      {:fname fname :lname lname :dob dob :phone phone})))

(defn connect-db
  "creates a db connection"
  [dbc]
  (open-db {:hostname (:host dbc)
            :port (:port dbc)
            :database (:dbname dbc)
            :username (:user dbc)
            :password (:password dbc)
            :pool-size 25}))

(defn close-db
  "closes a db connection"
  [db]
  (close-db! db))

(defn clj-add-person-db
  "add or update person to database's table"
  [db t p]
  ; I needed to add a primary key over fname lname and dob to make
  ; this work
  (let [ks (keys p)
        phone-k (name (last ks))]
    
```

```

      kcols (join " , " (map name (drop-last ks)))
      cols (join " , " (map name ks))
      vs (vals p)
      vals (str "'" (join "" , ' ' vs) "'") ;FIXME adding single quote works only if
      phone-v (str "'" (last vs) "'")
      query (str
        "INSERT INTO " t
        " (" cols ") "
        "VALUES (" vals ") "
        "ON CONFLICT (" kcols ") "
        "DO UPDATE SET " phone-k " = " phone-v "WHERE "phone-v" <> " t "." phone
      r (<!! (execute! db [query]))] ;FIXME we could extend the returning map
    (if (= {:updated 1 :rows []} r)
      {:updated 1 :ignored 0}
      {:updated 0 :ignored 1})))

(defn clj-add-people-db
  "add or update person to database's table"
  [db t ps]
  (let [f (fn [acc r] (let [r {:updated (+ (r :updated) (acc :updated))
                              :ignored (+ (r :ignored) (acc :ignored))}]
                        (print (str "\r" "Insertion stats: " r " of " (+ (r :ignored)
                                                                           r)))]
          (reduce f {:updated 0 :ignored 0} (map #(clj-add-person-db db t %) ps)))]

    (defn clj-del-person-db
      "delete person from database's table"
      [db t p]
      (<!! (execute! db [(str "DELETE FROM " t " WHERE fname='"(p :fname)'" AND lname='"(p

(defn xml-db-transfer-people
  "moves xml people records to database people rows"
  [db t url]
  (clj-add-people-db db t (xml->clj url)))

(defn -main
  "This converts the somecoolcompany huge xml in db data (it requires a postgres username
  [user]

```

```
(let [url "/tmp/somecoolcompany-backend-test/update-file.xml"
      dbc {:host "localhost"
            :port 5432
            :dbname "postgres"
            :user user
            :password ""}
      db (connect-db dbc)
      t "person"]
  (println (str "\n" "Finished:" (xml-db-transfer-people db t url)))
  (close-db db)))
```

6.7 DONE refactor code further ("Write clean code")

- CLOSING NOTE *[2017-01-22 Sun 19:55]*
I did refactoring before

```
(ns trycoolcompany-exercise.core
  (:require [clojure.zip :as zip]
             [clojure.data.zip.xml :refer :all]
             [clojure.data.xml :refer [parse]]
             [clojure.data.xml :as data.xml]
             [clojure.core.async :as async :refer [<!!]]
             [postgres.async :refer :all]
             [clojure.string :refer [join]]))

(defn xml->clj
  "transforms zipper nodes in maps"
  [url]
  (let [get-members (fn [x] (->> x
                                clojure.java.io/input-stream
                                parse
                                :content
                                (filter #(= :member (:tag %)))))
        get-value (fn [m k] (xml1-> (zip/xml-zip m) k text))]
    (for [member (get-members url)] :let [fname (get-value member :firstname)
                                           lname (get-value member :lastname)
                                           dob (get-value member :date-of-birth)
                                           phone (get-value member :phone)]
      ]
      {:fname fname :lname lname :dob dob :phone phone})))
```

```

(defn connect-db
  "creates a db connection"
  [dbc]
  (open-db {:hostname (:host dbc)
            :port (:port dbc)
            :database (:dbname dbc)
            :username (:user dbc)
            :password (:password dbc)
            :pool-size 25})))

(defn close-db
  "closes a db connection"
  [db]
  (close-db! db))

(defn clj-add-person-db
  "add or update person to database's table"
  [db t p]
  ; I needed to add a primary key over fname lname and dob to make
  ; this work
  (let [ks (keys p)
        phone-k (name (last ks))
        kcols (join " , " (map name (drop-last ks)))
        cols (join " , " (map name ks))
        vs (vals p)
        vals (str "'" (join "' , '" vs) "'") ;FIXME adding single quote works only if
        phone-v (str "'" (last vs) "'")
        query (str
                "INSERT INTO " t
                " (" cols ") "
                "VALUES (" vals ") "
                "ON CONFLICT (" kcols ") "
                "DO UPDATE SET " phone-k " = " phone-v "WHERE "phone-v" <> " t "." phone
                r (<!! (execute! db [query]))] ;FIXME we could extend the returning map
    (if (= {:updated 1 :rows []} r)
      {:updated 1 :ignored 0}
      {:updated 0 :ignored 1})))

```

```

(defn clj-add-people-db
  "add or update person to database's table"
  [db t ps]
  (let [f (fn [acc r] (let [r {:updated (+ (r :updated) (acc :updated))
                               :ignored (+ (r :ignored) (acc :ignored))}]
                        (print (str "\r" "Insertion stats: " r " of " (+ (r :ignored)
                                                                           r))))
        (reduce f {:updated 0 :ignored 0} (map #(clj-add-person-db db t %) ps)))]

    )

(defn clj-del-person-db
  "delete person from database's table"
  [db t p]
  (<!! (execute! db [(str "DELETE FROM " t " WHERE fname='"(p :fname)'" AND lname='"(p

(defn xml-db-transfer-people
  "moves xml people records to database people rows"
  [db t url]
  (clj-add-people-db db t (xml->clj url)))

(defn -main
  "This converts the somecoolcompany huge xml in db data (it requires a postgres username and password)"
  [user]
  (let [url "/tmp/somecoolcompany-backend-test/update-file.xml"
        dbc {:host "localhost"
              :port 5432
              :dbname "postgres"
              :user user
              :password ""}
        db (connect-db dbc)
        t "person"]
    (println (str "\n" "Finished:" (xml-db-transfer-people db t url)))
    (close-db db)))

```

6.8 DONE "minimize the overall run time of the merge process"

- CLOSING NOTE *[2017-01-19 Thu 17:49]*
 maybe I can increase the pool of connections? 25 does 1000 insertion every 5 seconds cc

- Note taken on *[2017-01-18 Wed 11:53]*

I think core.async should do a good job using threads. Maybe I can create more threads somehow? Check postgres.async doc.

<http://blog.korny.info/2014/03/08/xml-for-fun-and-profit.html>

With a pool of 100 inserting the xml data in an empty db (in memory)

runs in:

```
real 11m1.304s
user 7m28.874s
sys 0m52.740s
```

with a pool of 1500:

```
real 10m52.449s
user 7m27.015s
sys 0m52.778s
```

so it is irrelevant.

By the way tested on the real database:

```
Insertion stats: {:updated 927147, :ignored 572853} of 1500000
Statistics{:updated 927147, :ignored 572853}
```

```
real 9m4.016s
user 6m28.890s
sys 0m44.068s
```

6.9 DONE check performance (" Reason about performance and memory usage")

- CLOSING NOTE *[2017-01-22 Sun 20:25]*

constant usage of cpu

and time

```
Insertion stats: {:updated 927147, :ignored 572853} of 1500000
Statistics{:updated 927147, :ignored 572853}
```

```
real 9m4.016s user 6m28.890s sys 0m44.068s
```

Let's install sysstat

```
nix-env -iA nixos.sysstat
```

Lets get a performance stat of our process while it is running:

```
pidstat -p <pid-of-process-here> 2
```

Table 1:

time	UID	PID	%usr	%system	%guest	%CPU	CPU	Command
20:00:27	1000	15023	70.5	9.5	0.0	80.0	2	java
20:00:29	1000	15023	73.5	9.0	0.0	82.5	2	java
20:00:31	1000	15023	72.5	7.0	0.0	79.5	2	java
20:00:33	1000	15023	71.0	9.5	0.0	80.5	2	java
20:00:35	1000	15023	72.5	7.5	0.0	80.0	2	java
20:00:37	1000	15023	71.0	8.5	0.0	79.5	2	java
20:00:39	1000	15023	71.0	9.0	0.0	80.0	2	java
20:00:41	1000	15023	70.5	8.0	0.0	78.5	2	java
20:00:43	1000	15023	68.0	8.0	0.0	76.0	2	java
20:00:45	1000	15023	69.0	8.5	0.0	77.5	2	java
20:00:47	1000	15023	71.0	7.0	0.0	78.0	2	java
20:00:49	1000	15023	66.0	8.0	0.0	74.0	2	java
20:00:51	1000	15023	68.0	8.5	0.0	76.5	2	java
20:00:53	1000	15023	69.5	7.5	0.0	77.0	2	java
20:00:55	1000	15023	62.0	8.0	0.0	70.0	2	java
20:00:57	1000	15023	65.0	7.5	0.0	72.5	2	java
20:00:59	1000	15023	66.5	8.0	0.0	74.5	2	java
20:01:01	1000	15023	68.0	8.0	0.0	76.0	2	java
20:01:03	1000	15023	67.5	8.5	0.0	76.0	2	java
20:01:05	1000	15023	67.0	9.0	0.0	76.0	2	java
20:01:07	1000	15023	68.0	8.0	0.0	76.0	2	java
20:01:09	1000	15023	66.5	7.5	0.0	74.0	2	java
20:01:11	1000	15023	70.0	7.5	0.0	77.5	2	java
20:01:13	1000	15023	69.0	9.0	0.0	78.0	2	java
20:01:15	1000	15023	70.5	6.5	0.0	77.0	2	java
20:01:17	1000	15023	70.0	8.0	0.0	78.0	2	java
20:01:19	1000	15023	62.0	8.5	0.0	70.5	2	java
20:01:21	1000	15023	66.0	7.0	0.0	73.0	2	java
20:01:23	1000	15023	68.0	8.0	0.0	76.0	2	java
20:01:25	1000	15023	68.5	8.0	0.0	76.5	2	java
20:01:27	1000	15023	68.5	8.0	0.0	76.5	2	java
20:01:29	1000	15023	69.0	7.5	0.0	76.5	2	java
20:01:31	1000	15023	68.5	8.0	0.0	76.5	2	java
20:01:33	1000	15023	67.5	7.5	0.0	75.0	2	java
20:01:35	1000	15023	68.0	8.0	0.0	76.0	2	java
20:01:37	1000	15023	69.0	7.5	0.0	76.5	2	java
20:01:39	1000	15023	63.5	8.0	0.0	71.5	2	java
20:01:41	1000	15023	69.0	8.0	0.0	77.0	2	java
20:01:43	1000	15023	67.0	8.0	0.0	75.0	2	java
20:01:45	1000	15023	67.5	9.0	0.0	76.5	2	java
20:01:47	1000	15023	68.5	6.5	0.0	75.0	2	java
20:01:49	1000	15023	67.0	9.0	0.0	76.0	2	java
20:01:51	1000	15023	68.5	8.0	0.0	76.5	2	java
20:01:53	1000	15023	69.0	6.5	0.0	75.5	2	java
20:01:55	1000	15023	69.0	8.0	0.0	77.0	2	java
20:01:57	1000	15023	69.0	7.5	0.0	76.5	2	java
20:01:59	1000	15023	69.5	6.5	0.0	76.0	2	java

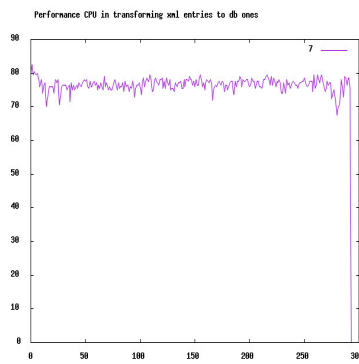


Figure 1: Performance in inserting xml entries in db